

Assignment: Practicing Ensemble Learning in Python

Objective:

Learn to implement various ensemble methods and understand their differences through practice on real datasets.

Step 0: Setup

Import them in your notebook:

```
Python
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split,
cross_val_score

from sklearn.metrics import accuracy_score,
classification_report, confusion_matrix

from sklearn.tree import DecisionTreeClassifier

from sklearn.ensemble import RandomForestClassifier,
AdaBoostClassifier, GradientBoostingClassifier,
VotingClassifier, BaggingClassifier

from sklearn.linear_model import LogisticRegression

from sklearn.svm import SVC
```



Step 1: Load Dataset

Choose one of the following datasets:

1. Heart Disease - [Kaggle Link](#)
2. Diabetes - [Kaggle Link](#)

Python

```
df = pd.read_csv("path_to_dataset.csv")  
  
df.head()
```

Tasks:

- Check for missing values
- Explore the data (e.g., `df.describe()`, `df.info()`)

Step 2: Preprocessing

- Split into features and target variable:

Python

```
X = df.drop('target_column', axis=1)  
  
y = df['target_column']
```

Split into training and testing sets:

Python

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```



Step 3: Practice Bagging & Pasting

1. Train a Decision Tree with Bagging:

Python

```
bagging_model = BaggingClassifier(  
  
    base_estimator=DecisionTreeClassifier(),  
  
    n_estimators=10,  
  
    max_samples=0.8, # 80% samples  
  
    bootstrap=True, # Bagging (with replacement)  
  
    random_state=42  
  
)  
  
bagging_model.fit(X_train, y_train)  
  
y_pred = bagging_model.predict(X_test)  
  
print("Bagging Accuracy:", accuracy_score(y_test, y_pred))
```

2. Change `bootstrap=False` to practice Pasting. Compare accuracy.



Step 4: Practice Random Forest

Python

```
rf_model = RandomForestClassifier(n_estimators=50, random_state=42)

rf_model.fit(X_train, y_train)

y_pred = rf_model.predict(X_test)

print("Random Forest Accuracy:", accuracy_score(y_test, y_pred))
```

- Plot feature importance:

Python

```
importances = rf_model.feature_importances_

sns.barplot(x=importances, y=X.columns)

plt.show()
```

Step 5: Practice Boosting

1. AdaBoost:

Python

```
ada_model = AdaBoostClassifier(

    base_estimator=DecisionTreeClassifier(max_depth=1),

    n_estimators=50,

    learning_rate=1,

    random_state=42

)
```



```
ada_model.fit(X_train, y_train)

y_pred = ada_model.predict(X_test)

print("AdaBoost Accuracy:", accuracy_score(y_test, y_pred))
```

2. Gradient Boosting:

Python

```
gb_model = GradientBoostingClassifier(n_estimators=50, learning_rate=0.1,
random_state=42)

gb_model.fit(X_train, y_train)

y_pred = gb_model.predict(X_test)

print("Gradient Boosting Accuracy:", accuracy_score(y_test, y_pred))
```

Step 6: Practice Voting

Python

```
voting_model = VotingClassifier(
    estimators=[

        ('lr', LogisticRegression(max_iter=1000)),

        ('dt', DecisionTreeClassifier()),

        ('rf', RandomForestClassifier(n_estimators=50))
    ],
    voting='hard' # try 'soft' as well
)
```



```
voting_model.fit(X_train, y_train)

y_pred = voting_model.predict(X_test)

print("Voting Accuracy:", accuracy_score(y_test, y_pred))
```

Step 7: Practice Stacking

Python

```
from sklearn.ensemble import StackingClassifier

stack_model = StackingClassifier(
    estimators=[

        ('lr', LogisticRegression(max_iter=1000)),

        ('rf', RandomForestClassifier(n_estimators=50)),

        ('dt', DecisionTreeClassifier())
    ],
    final_estimator=GradientBoostingClassifier(n_estimators=50),
)

stack_model.fit(X_train, y_train)

y_pred = stack_model.predict(X_test)

print("Stacking Accuracy:", accuracy_score(y_test, y_pred))
```



Step 8: Analysis & Report

1. Compare all models' accuracies.
2. Discuss:
 - Which model performed best?
 - Difference between Bagging and Boosting?
 - Advantages of Stacking?
3. Optional: Plot confusion matrices for each model.

Deliverables:

- Python notebook with all code, plots, and outputs.
- Short written summary comparing the models.

